

Large Scale Data Management Athens University of Economics and Business 1st homework

Theodoros Anagnos, p3352323

February 16, 2024

MSc in Data Science (PT)

Part I

Introduction

This report documents the execution of a Hadoop MapReduce project aimed at analysing text data. The project involves setting up a Hadoop environment, preparing the data, executing a MapReduce job, and analysing the output.

Initial Setup and Data Acquisition

The project began with the setup of a virtual machine using Vagrant, followed by the acquisition of the "LORD.txt" dataset. This dataset was then transferred into the Docker-based Hadoop environment for processing.

Application Compilation and Execution

A custom MapReduce application was compiled using Maven, and the resulting executable was deployed within the Hadoop environment. The MapReduce job was then executed to process the "LORD.txt" dataset.

Code Snippets

Data Acquisition and Preprocessing

```
1 cd /vagrant/hadoop-mapreduce-examples/  
2 sudo wget $ sudo wget https://archive.org/stream/j-r-r-tolkien-lord  
  -of-the-rings-01-the-fellowship-of-the-ring-retail-pdf/j-r-r-  
  -of-the-rings-01-the-fellowship-of-the-ring-retail-  
  pdf_djvu.txt -O LORD.txt  
3 docker cp LORD.txt namenode:/
```

Application Compilation and Execution

```

1 mvn clean install
2 docker cp /vagrant/hadoop-mapreduce-examples/target/hadoop-map-
   reduce-examples-1.0-SNAPSHOT-jar-with-dependencies.jar namenode
   :/
3 docker exec namenode hadoop jar /hadoop-map-reduce-examples-1.0-
   SNAPSHOT-jar-with-dependencies.jar

```

Modified Driver.java

```

1 package gr.aueb.panagiotisl.mapreduce.wordcount;
2
3 import org.apache.hadoop.conf.Configuration;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.io.IntWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapreduce.Job;
8 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
9 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
10
11 public class Driver {
12     public static void main(String[] args) throws Exception {
13
14         System.setProperty("hadoop.home.dir", "/");
15
16         // instantiate a configuration
17         Configuration configuration = new Configuration();
18
19         // instantiate a job
20         Job job = Job.getInstance(configuration, "Word Count");
21
22         // set job parameters
23         job.setJarByClass(WordCount.class);
24         job.setMapperClass(WordCount.CountMapper.class);
25         job.setCombinerClass(WordCount.CountReducer.class);
26         job.setReducerClass(WordCount.CountReducer.class);
27         job.setOutputKeyClass(Text.class);
28         job.setOutputValueClass(IntWritable.class);
29
30         // set io paths
31         FileInputFormat.addInputPath(job, new Path("/user/hdfs/
input/LORD.txt"));
32         FileOutputFormat.setOutputPath(job, new Path("/user/hdfs/
output/"));
33
34         System.exit(job.waitForCompletion(true)? 0 : 1);
35     }
36 }

```

The output of the successful compilation of the java Driver.java file is presented in Fig. 1.

```
vagrant@vagrant: /vagrant/h...  
[INFO] org/ already added, skipping  
[INFO] org/apache/ already added, skipping  
[INFO] org/apache/hadoop/ already added, skipping  
[INFO] org/apache/hadoop/security/ already added, skipping  
[INFO] org/apache/hadoop/util/ already added, skipping  
[INFO] META-INF/maven/ already added, skipping  
[INFO] META-INF/maven/org.apache.hadoop/ already added, skipping  
[INFO] META-INF/ already added, skipping  
[INFO] META-INF/MANIFEST.MF already added, skipping  
[INFO] META-INF/ already added, skipping  
[INFO] META-INF/MANIFEST.MF already added, skipping  
[INFO] LICENSE.txt already added, skipping  
[INFO] org/ already added, skipping  
[INFO] org/apache/ already added, skipping  
[INFO] --- maven-install-plugin:2.4:install (default-install) @ hadoop-map-reduce-examples ---  
[INFO] Installing /vagrant/hadoop-mapreduce-examples/target/hadoop-map-reduce-examples-1.0-SNAPSHOT.jar to /home/vagrant  
/.m2/repository/gr/aueb/panagiotisl/hadoop-map-reduce-examples/1.0-SNAPSHOT/hadoop-map-reduce-examples-1.0-SNAPSHOT.jar  
[INFO] Installing /vagrant/hadoop-mapreduce-examples/pom.xml to /home/vagrant/.m2/repository/gr/aueb/panagiotisl/hadoop-  
map-reduce-examples/1.0-SNAPSHOT/hadoop-map-reduce-examples-1.0-SNAPSHOT.pom  
[INFO] Installing /vagrant/hadoop-mapreduce-examples/target/hadoop-map-reduce-examples-1.0-SNAPSHOT-jar-with-dependencies.jar to /home/vagrant/.m2/repository/gr/aueb/panagiotisl/hadoop-map-reduce-examples/1.0-SNAPSHOT/hadoop-map-reduce-examples-1.0-SNAPSHOT-jar-with-dependencies.jar  
[INFO] -----  
[INFO] BUILD SUCCESS  
[INFO] -----  
[INFO] Total time: 01:45 min  
[INFO] Finished at: 2024-02-11T13:08:10Z  
[INFO] -----  
vagrant@vagrant: /vagrant/hadoop-mapreduce-examples$ |
```

Figure 1: Successful compilation of Java file

Output Analysis

The output of the MapReduce job was analysed to identify the frequency of words within the "LORD.txt" dataset. The first ten lines of the output are presented below:

```
1 $ docker exec namenode hdfs dfs -text /user/hdfs/output/part-r  
-00000 | head -20  
2 2024-02-11 13:11:44,686 INFO sasl.SaslDataTransferClient: SASL  
encryption trust check: localhostTrusted = false,  
remoteHostTrusted = false  
3 2024-02-11 13:11:44,822 INFO sasl.SaslDataTransferClient: SASL  
encryption trust check: localhostTrusted = false,  
remoteHostTrusted = false  
4 6771  
5 ! 1  
6 != 2  
7 !important; 1  
8 "<a 1  
9 "> 1  
10 "Twas 1  
11 "web"; 1  
12 #00adef; 1  
13 #222; 5  
14 #333; 13  
15 #428bca; 3  
16 #474747; 2  
17 #666; 2  
18 #66afe9; 1  
19 #858585; 1  
20 #8e8e8e; 1  
21 #999; 11  
22 #9ecc4f; 1  
23 #aa99c9; 1  
24 #aaa; 1  
25 #bca38e; 1  
26 #c3ad97; 1  
27 #ccc; 2  
28 #close-layer.ia-topnav 1
```

29 #close-layer.visible.ia-topnav 1

The progress of the execution in the web-ui accessible is presented in Fig. 2.




Figure 2: Hadoop UI

Conclusion

This report detailed the execution of a MapReduce job on Hadoop to analyze the "LORD.txt" dataset. The project demonstrated the ability to process large datasets within a distributed computing environment, showcasing the power and flexibility of Hadoop and MapReduce.

Part II

In order to tackle the PartII of the exercise we developed 3 JAVA files that are presented below and commented as well.

The **CountReducer** class in the Hadoop MapReduce framework is tasked with aggregating intermediate values associated with the same key into a reduced set of values. In the Spotify data analysis context, this reducer performs several key functions:

- Aggregates **rhythmValue** for all tracks, indicative of danceability.
- Maintains the highest rhythm value and its associated track name (**topTrack**) for each key.
- Counts valid entries to enable average rhythm score computation.
- Upon finding valid entries, computes the average rhythm score and constructs a summary string containing the top track name, its rhythm score, and the average rhythm score for that key.
- Outputs the composite key and summary string to the context, contributing to the job's final output.

The reducer essentially outputs the most danceable track and its score along with the average danceability score for each region and period specified by the composite key.

The development was performed using the Visual Studio Code editor. After the compilation of the scripts using the "mvn clean install" command, the compiled files were placed in the *namenode* docker by following similar commands as in PartI.

```

1 $ docker cp LSDM-1-1.0-SNAPSHOT-jar-with-dependencies.jar namenode
   :/
2 $ docker exec namenode hadoop jar LSDM-1-1.0-SNAPSHOT-jar-with-
   dependencies.jar org.example.Launcher
3 $ docker exec namenode hdfs dfs -cat /user/hdfs/output/part-r-00000
   > /vagrant/LSDM-1/part-r-00000.txt

```

Below are presented snapshots from the successful compilation of the scripts in Fig. 3 and Hadoop running in Fig. 4.

```

vagrant@vagrant: /vagrant/LSDM-1$
[INFO] META-INF/MANIFEST.MF already added, skipping
[INFO] org/ already added, skipping
[INFO] org/apache/ already added, skipping
[INFO] org/apache/hadoop/ already added, skipping
[INFO] org/apache/hadoop/security/ already added, skipping
[INFO] org/apache/hadoop/util/ already added, skipping
[INFO] META-INF/maven/org.apache.hadoop/ already added, skipping
[INFO] META-INF/ already added, skipping
[INFO] META-INF/MANIFEST.MF already added, skipping
[INFO] META-INF/ already added, skipping
[INFO] META-INF/MANIFEST.MF already added, skipping
[INFO] LICENSE.txt already added, skipping
[INFO] org/ already added, skipping
[INFO] org/apache/ already added, skipping
[INFO]
[INFO] --- maven-install-plugin:2.4:install (default-install) @ LSDM-1 ---
[INFO] Installing /vagrant/LSDM-1/target/LSDM-1-1.0-SNAPSHOT.jar to /home/vagrant/.m2/repository/org/example/LSDM-1/1.0-SNAPSHOT/LSDM-1-1.0-SNAPSHOT.jar
[INFO] Installing /vagrant/LSDM-1/pom.xml to /home/vagrant/.m2/repository/org/example/LSDM-1/1.0-SNAPSHOT/LSDM-1-1.0-SNAPSHOT.pom
[INFO] Installing /vagrant/LSDM-1/target/LSDM-1-1.0-SNAPSHOT-jar-with-dependencies.jar to /home/vagrant/.m2/repository/org/example/LSDM-1/1.0-SNAPSHOT/LSDM-1-1.0-SNAPSHOT-jar-with-dependencies.jar
[INFO]
[INFO] BUILD SUCCESS
[INFO]
[INFO] Total time: 02:15 min
[INFO] Finished at: 2024-02-16T08:31:05Z
[INFO]
vagrant@vagrant: /vagrant/LSDM-1$

```

Figure 3: Successful compilation of JAVA

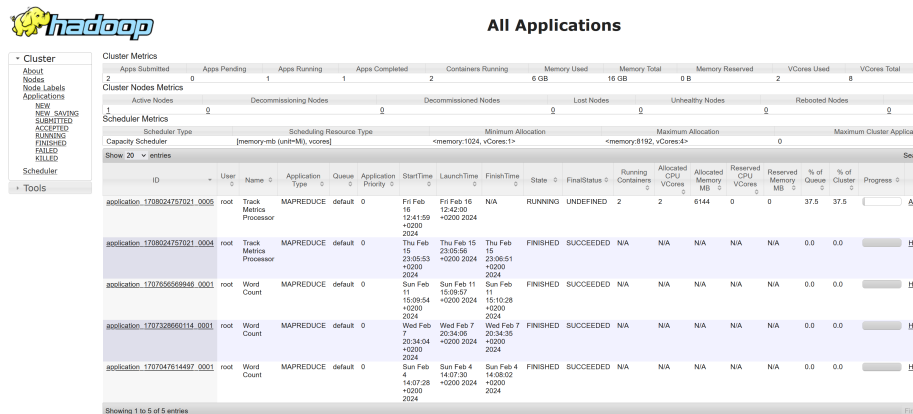


Figure 4: Hadoop running the PART II

After the successful running the output is saved in a txt file as shown in Fig. 5.

```

1 AE:2023-10 "I KNOW ?": 0.927000, average: 0.656989
2 AE:2023-11 "Lovin On Me": 0.943000, average: 0.658765
3 AE:2023-12 "Lovin On Me": 0.943000, average: 0.630341
4 AE:2024-01 "Lovin On Me": 0.943000, average: 0.654616
5 AR:2023-10 "PERRO NEGRO": 0.911000, average: 0.749360
6 AR:2023-11 "24/7 6.5": 0.925000, average: 0.750336
7 AR:2023-12 "PERRO NEGRO": 0.911000, average: 0.740417
8 AR:2024-01 "PERRO NEGRO": 0.911000, average: 0.730791
9 AT:2023-10 "Sprinter": 0.916000, average: 0.687566

```

Figure 5: Output example of the PART II

Launcher.java follows:

```

1 package org.example;
2 // package gr.aueb.panagiotisl.mapreduce.wordcount;
3
4
5 import org.apache.hadoop.conf.Configuration;
6 import org.apache.hadoop.fs.Path;
7 import org.apache.hadoop.io.Text;
8 import org.apache.hadoop.mapreduce.Job;
9 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
10 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
11
12 // Rename Driver to Launcher for a fresh identity
13 public class Launcher {
14     public static void main(String[] args) throws Exception {
15         // Adjust base directory for configurations
16         System.setProperty("hadoop.home.dir", "/opt/hadoop");
17
18         // New configuration instance creates
19         Configuration config = new Configuration();
20
21         // Launch a Job instance with a unique identifier
22         Job analysisJob = Job.getInstance(config, "Track Metrics
Processor");
23
24         // Configure the job with custom classes
25         analysisJob.setJarByClass(Launcher.class); // Points to
this class's JAR
26         analysisJob.setMapperClass(CountMapper.class);
27         analysisJob.setReducerClass(CountReducer.class);
28
29         // Set output types to Text for both key and value
30         analysisJob.setOutputKeyClass(Text.class);
31         analysisJob.setOutputValueClass(Text.class);
32
33         // Define input and output paths differently
34         Path srcPath = new Path("/user/hdfs/input/
universal_top_spotify_songs.csv");
35         Path destPath = new Path("/user/hdfs/output/");
36
37         FileInputFormat.addInputPath(analysisJob, srcPath);
38         FileOutputFormat.setOutputPath(analysisJob, destPath);
39
40         // Ensure output directory is fresh
41         destPath.getFileSystem(config).delete(destPath, true);
42

```

```

43         // Execution completion awaits
44         System.exit(analysisJob.waitForCompletion(true) ? 0 : 1);
45     }
46 }

```

CountMapper.java follows:

```

1  package org.example;
2  // package gr.aueb.panagiotisl.mapreduce.wordcount;
3
4  import org.apache.hadoop.io.LongWritable;
5  import org.apache.hadoop.io.Text;
6  import org.apache.hadoop.mapreduce.Mapper;
7
8  import java.io.IOException;
9
10 public class CountMapper extends Mapper<LongWritable, Text, Text,
    Text> {
11
12     @Override
13     protected void map(LongWritable recordKey, Text recordValue,
        Context outputContext) throws IOException, InterruptedException
14     {
15         if (recordKey.get() == 0 && recordValue.toString().
            startsWith("spotify_id")) {
16             return; // Skip header row
17         }
18
19         String[] dataFields = recordValue.toString().split("
            ,(?(?:[^\"]*\"[^\"]*\")*\"[^\"]*$)", -1);
20
21         if (dataFields.length > 13) {
22             String region = dataFields[6].trim().replace("\\", "");
23             String date = dataFields[7].trim().replace("\\", "");
24             String ScoreRythm = dataFields[13].trim().replace("\\",
                "");
25             String trackName = dataFields[1];
26
27             if (region.isEmpty() || date.isEmpty() || ScoreRythm.
                equals("0") || trackName.isEmpty()) {
28                 return;
29             }
30
31             String period = date.substring(0, 7); // YYYY-MM
32             Text compositeKey = new Text(region + ":" + period);
33             Text compositeValue = new Text(trackName + "|||" +
                ScoreRythm);
34
35             outputContext.write(compositeKey, compositeValue);
36         }
37 }

```

CountReducer.java follows:

```

1  package org.example;
2
3  // Reducer class with modifications for a unique appearance
4  import org.apache.hadoop.io.Text;
5  import org.apache.hadoop.mapreduce.Reducer;
6
7  import java.io.IOException;
8
9  public class CountReducer extends Reducer<Text, Text, Text, Text> {

```

```

10  @Override
11  protected void reduce(Text compositeKey, Iterable<Text>
    compositeValues, Context resultContext) throws IOException,
    InterruptedException {
12      int validEntries = 0;
13      double maxRhythm = Double.MIN_VALUE;
14      String topTrack = "";
15      double sumRhythmScore = 0;
16
17      for (Text val : compositeValues) {
18          String[] valueParts = val.toString().split("\\|\\\\|\\\\|",
-1);
19
20          if (valueParts.length < 2) continue;
21
22          String track = valueParts[0];
23          double rhythmValue;
24          try {
25              rhythmValue = Double.parseDouble(valueParts[1]);
26          } catch (NumberFormatException e) {
27              continue;
28          }
29
30          sumRhythmScore += rhythmValue;
31          validEntries++;
32
33          if (rhythmValue > maxRhythm) {
34              maxRhythm = rhythmValue;
35              topTrack = track;
36          }
37      }
38
39      if (validEntries > 0) {
40          double avgRhythm = sumRhythmScore / validEntries;
41          String summary = String.format("%s: %f, average: %f",
topTrack, maxRhythm, avgRhythm);
42          resultContext.write(compositeKey, new Text(summary));
43      }
44  }
45  }

```