


ATHENS
UNIVERSITY
OF
ECONOMICS &
BUSINESS

M.Sc. in Data Science



Text Analytics

Instructor: Ion Androutsopoulos

Assignment 04 - Natural Language Processing with Recurrent Neural Networks

Anagnos Theodoros (p3352323) - Michalopoulos Ioannis
(p3352314) - Kafantaris Panagiotis (p3352328) - Vigkos Ioannis
(p3352326)

30.05.2024

Contents

1	Exercise 01	2
2	Dataset and Preprocessing	2
2.1	Dataset	2
2.2	Preprocessing	3
3	Baseline Classifiers	4
3.1	Logistic Regression Classifier	4
4	RNN Model	4
4.1	Text Vectorization and Embedding Matrix Preparation	4
4.2	Self-Attention Layer	4
4.3	Build, Prepare Data for Training and Compile the Model	4
5	Hyperparameter Tuning	5
6	Results	5
6.1	Learning Curves	5
6.2	Metrics for Model Evaluation	7
6.2.1	Baseline	7
6.2.2	RNN Model	7
7	Conclusion	7
8	Exercise 02	8
8.1	Introduction	8
8.2	Dataset and Preprocessing	8
8.3	Dataset	8
8.4	Preprocessing	8
8.5	Baseline Models	8
8.6	Most Frequent Class (MFC) Baseline	8
8.7	MLP Baseline	9
8.8	MLP Model Summary	10
8.9	Bi-Directional Stacked RNN	13
8.10	RNN Model Summary	13
8.11	Architecture	14
8.12	Training and Hyperparameter Tuning	15
8.13	Results	16
8.14	Evaluation Metrics	16
8.15	Precision-Recall AUC for Each Class	17
8.16	Insights from Precision-Recall AUC	19
8.17	Insights and Discussion	19
8.18	Hyperparameter Tuning Insights	19
8.19	Impact of Data Preprocessing	20
8.20	Conclusion	20

1 Exercise 01

In this report, we

The implementation details and code can be accessed through the following Colab notebook link:

[Link here](#)

2 Dataset and Preprocessing

2.1 Dataset

We used a Twitter dataset from Kaggle for our sentiment analysis task. The dataset includes tweets labeled with sentiment categories: positive, negative, and neutral. The dataset statistics are as follows:

- Number of records: *4870*
- Number of classes: 3 (Positive, Negative, Neutral)
- Vocabulary size: *6000 terms, vectorized with TF-IDF*

Provenance:

- This dataset enriches a benchmark dataset available at: <https://mcrlab.net/research/mvsa-sentiment-analysis-on-multi-view-social-data/>

Collection Methodology:

- This dataset is enriched by augmenting a dataset of images and captions with sentiment analysis-generated labels. A method was put forth to determine the ideal sentiment using the caption feature to implement six different sentiment analysis techniques. Final classification of each dataset item is based on the percentage of positive, negative, and neutral polarities extracted for each caption.

Sample:

LabeledText .xlsx

File Edit View Insert Format Data Tools Help

Menu 100% \$ % .0 .00 123 Calibri - [14] + B I ÷ A

A1 File Name

	A	B	C
1	File Name	Caption	LABEL
2	1.txt	How I feel today #legday #jelly #aching #gym	negative
3	10.txt	@ArrivaTW absolute disgrace two carriages from Bangor half way there standing room only #disgraced	negative
4	100.txt	This is my Valentine's from 1 of my nephews. I am elated; sometimes the little things are the biggest & best things!	positive
5	1000.txt	betterfeelingfilms: RT via Instagram: First day of filming #powerless back in 2011. Can't i	neutral
6	1001.txt	Zoe's first love #Rattled @JohnnyHarper15	positive
7	1002.txt	Chaotic Love - giclee print ?65 at #art #love #chaotic #abstract #blue #silver #prints #buy	positive
8	1003.txt	They gna be mad when I reach that goal though. #Rejected the wrong girl ? just getting started & already turn heads.?	negative
9	1004.txt	On day 9.. It's now in my daily routine.. Feeling guuuurrrddd! ? #Aching #PainNoGain #FeelingGood	negative
10	1005.txt	#ANIMALABUSE #TORONTO #PUDDY #TORTURE WE OFFER \$1K #REWARD puppy #beaten #bound #burned	neutral
11	1006.txt	Mike will not accept this plastic rose. @wfaamike @wfaachannel8 @wfaagmt #rejected	negative
12	1007.txt	Just ate four cookies. #remorse	negative
13	1008.txt	It's shocking what is acceptable in kids TV shows these days #shocking #shocked	negative
14	1009.txt	We are so #excited to announce that we have launched our #affiliate program please visit us at	positive
15	101.txt	Just when you thought you'd seen everything! #ParkingSpace #Lanzarote #ItsNOTtaxed #zimmer #oldpeople #alarmed ?	negative
16	1010.txt	RT @MissGem: So this @parcelforce van thinks it's ok to nearly run people off the road?! #disgusting #shocked #disgraceful	negative
17	1011.txt	Today I #StepBackInTime !!! @PWLHitFactory @kylieminogue #BetterTheDevilYouKnow #WhatDolHaveToDo #Shocked	neutral
18	1012.txt	Photos: #Photographer got a rumble in the #jungle as he was #beaten by 30-stone #gorilla	neutral

Figure 1: Dataset sample

2.2 Preprocessing

The preprocessing steps included:

- **Converting text to lowercase:** This step helps in normalizing the text data by removing case sensitivity, ensuring that words like "Happy" and "happy" are treated the same.
- **Removing URLs, HTML tags, punctuation, numbers, and extra spaces:** These elements are generally not useful for sentiment analysis and can be considered as noise. Removing them helps in focusing on the actual textual content.
- **Converting chat words to their full forms:** Chat words (e.g., "idk" to "I don't know") are expanded to their full forms to enhance the understanding of the text.

- **Tokenization and removing stop words:** Tokenization splits the text into individual words, and stop words (common words like "the", "is") are removed to reduce noise and focus on the significant words.
- **Lemmatization using NLTK:** Lemmatization reduces words to their base or root form (e.g., "running" to "run"), which helps in treating different forms of a word as a single item.

3 Baseline Classifiers

3.1 Logistic Regression Classifier

We used TF-IDF (Term Frequency-Inverse Document Frequency) features with unigram and bigram representations to capture both individual words and pairs of consecutive words. TF-IDF helps in identifying the importance of words in a document relative to the entire dataset.

Logistic Regression is a linear model commonly used for binary and multi-class classification tasks. It estimates the probability that a given input belongs to a particular class by applying the logistic function. In our case, we trained a logistic regression classifier on TF-IDF features to serve as a baseline for comparison with more complex models.

4 RNN Model

4.1 Text Vectorization and Embedding Matrix Preparation

We preprocess text data by converting it into sequences of integers using the TextVectorization layer. It then prepares an embedding matrix using pre-trained word2vec embeddings, mapping each word in the vocabulary to its corresponding vector. This embedding matrix can later be used in the model to convert integer sequences back into dense vector representations.

4.2 Self-Attention Layer

The SelfAttention class implements a self-attention mechanism that computes attention scores for each time step in an input sequence, normalizes them, and uses them to create a weighted sum of the input features. This process helps the model focus on the most relevant parts of the input sequence when making predictions. The class is designed to be used as a layer in a Keras model, allowing it to be easily integrated into larger neural network architectures.

4.3 Build, Prepare Data for Training and Compile the Model

We construct a neural network for text classification using TensorFlow and Keras. It includes layers for text vectorization, embedding with pre-trained word vectors, bidirectional LSTM for sequence modeling, self-attention for focusing on important parts of the input, and dense layers for classification. The model

is compiled, data is prepared, and the model is trained with a validation set to monitor performance.

5 Hyperparameter Tuning

We perform hyperparameter tuning for a neural network model using the Keras Tuner library. It defines a model-building function, initializing the tuner, and running the search to find the optimal hyperparameters. Trains the model using these optimal settings. The process aims to improve model performance by systematically exploring different configurations.

- The optimal number of units in the first densely-connected layer is 50.
- The optimal LSTM size is 200.
- The optimal number of MLP layers is 0.
- The optimal MLP units is 50.
- The optimal dropout rate is 0.2.

6 Results

6.1 Learning Curves

Figure 2 and Figure 3 present the training and validation accuracy and loss over 8 epochs. In Figure 2, the training accuracy shows a steady increase, reaching approximately 0.8, while the validation accuracy improves initially but then fluctuates around 0.7. This indicates that the model is learning, but its performance on validation data is not consistently improving. Figure 3 shows the loss curves, where the training loss decreases continuously, signifying effective learning on training data. However, the validation loss initially decreases but then remains relatively flat, suggesting that the model's ability to generalize to new data does not improve significantly after the first few epochs.




Figure 2: Learning curves showing the model accuracy on the training and development sets as a function of epochs.




Figure 3: Learning curves showing the model loss on the training and development sets as a function of epochs.

6.2 Metrics for Model Evaluation

6.2.1 Baseline

Table 1: Performance Metrics for Logistic Regression on the Test Subset

Class	Precision	Recall	F1-Score	Support
Negative	0.70	0.64	0.67	221
Neutral	0.61	0.66	0.63	273
Positive	0.74	0.73	0.73	237
Accuracy	0.67 (731)			
Macro Avg	0.68	0.67	0.68	731
Weighted Avg	0.68	0.67	0.68	731

6.2.2 RNN Model

Table 2: Performance Metrics for Logistic Regression on the Test Subset

Class	Precision	Recall	F1-Score	Support
Negative	0.74	0.67	0.70	221
Neutral	0.69	0.63	0.66	273
Positive	0.73	0.88	0.80	237
Accuracy	0.72 (731)			
Macro Avg	0.72	0.72	0.72	731
Weighted Avg	0.72	0.72	0.72	731

7 Conclusion

The neural network model shows a clear improvement over the logistic regression model in terms of accuracy, precision, recall, and F1-score. The significant improvement in recall and F1-score for the positive class and the overall higher macro and weighted averages indicate that the neural network model is better at handling the complexities of the text classification task.

8 Exercise 02

8.1 Introduction

In this report, we repeat exercise 10 of Part 3 by developing a bi-directional stacked RNN (with LSTM cells) for POS tagging. This task involves using Keras/TensorFlow to implement the model, tuning the hyperparameters, and comparing the performance with baseline models.

The implementation details and code can be accessed through the following Colab notebook link: [Link here](#)

8.2 Dataset and Preprocessing

8.3 Dataset

We used the Universal Dependencies English Web Treebank dataset for our POS tagging task. The dataset includes sentences annotated with POS tags. The statistics are as follows:

- Number of sentences: 12,543 (train), 2,002 (dev), 2,077 (test)
- Vocabulary size: 34,078 terms (vectorized with Word2Vec)

Provenance:

- Universal Dependencies: <http://universaldependencies.org/>

8.4 Preprocessing

The preprocessing steps included:

- Converting text to lowercase.
- Removing punctuation, numbers, and extra spaces.
- Creating sliding windows for words.
- Mapping words to pre-trained Word2Vec embeddings.
- Tokenizing windows using the created vocabulary.
- Encoding POS tags using LabelEncoder and OneHotEncoder.

8.5 Baseline Models

8.6 Most Frequent Class (MFC) Baseline

The MFC baseline tags each word with the most frequent tag from the training data. This simple approach provides a benchmark for evaluating more complex models.

Table 3: Performance Metrics for MFC Baseline on the Test Subset

Class	Precision	Recall	F1-Score	Support
NOUN	0.16	1.00	0.28	2,077
ADJ	0.00	0.00	0.00	0
...
Accuracy	0.16 (2,077)			
Macro Avg	0.04	0.14	0.05	2,077
Weighted Avg	0.03	0.16	0.05	2,077

The most frequent tag in the training data is: NOUN
Baseline Model Accuracy: 0.1625540275049116

Classification Report for the Baseline Model:

	precision	recall	f1-score	support
ADJ	0.00	0.00	0.00	1788
ADP	0.00	0.00	0.00	2034
ADV	0.00	0.00	0.00	1176
AUX	0.00	0.00	0.00	1543
CCONJ	0.00	0.00	0.00	737
DET	0.00	0.00	0.00	1897
INTJ	0.00	0.00	0.00	120
NOUN	0.16	1.00	0.28	4137
NUM	0.00	0.00	0.00	542
PART	0.00	0.00	0.00	649
PRON	0.00	0.00	0.00	2162
PROPN	0.00	0.00	0.00	2077
PUNCT	0.00	0.00	0.00	3096
SCONJ	0.00	0.00	0.00	384
SYM	0.00	0.00	0.00	109
VERB	0.00	0.00	0.00	2606
X	0.00	0.00	0.00	39
-	0.00	0.00	0.00	354
accuracy			0.16	25450
macro avg	0.01	0.06	0.02	25450
weighted avg	0.03	0.16	0.05	25450

Figure 4: Performance Metrics for MFC Baseline on the Test Subset

8.7 MLP Baseline

We implemented an MLP baseline using Keras. The architecture included an embedding layer, two dense layers with ReLU activation, and dropout for regularization. The model was trained using categorical cross-entropy loss and the Adam optimizer.

8.8 MLP Model Summary

The following table shows the architecture of the MLP model. It includes an embedding layer to convert words into dense vectors, followed by flatten, dense, and dropout layers. The model uses a softmax activation in the output layer for multi-class classification.

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 5, 300)	5,965,200
flatten (Flatten)	(None, 1500)	0
dense (Dense)	(None, 256)	384,256
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32,896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 18)	2,322

Total params: 6,384,674 (24.36 MB)

Trainable params: 6,384,674 (24.36 MB)

Non-trainable params: 0 (0.00 B)

Figure 5: MLP Model Summary

Figure 6 and Figure 7 depict the adjusted MLP model's accuracy and loss over 25 epochs for both training and validation datasets. In Figure 6, the training accuracy shows a rapid increase, stabilizing around 0.90, while the validation accuracy closely follows, indicating good generalization. Figure 7 shows the loss curves, where the training loss decreases significantly, leveling off below 0.8, while the validation loss initially decreases and then fluctuates around 0.8. The close alignment of training and validation metrics suggests that the adjustments made to the MLP model have improved its performance, reducing overfitting and enhancing generalization to validation data.

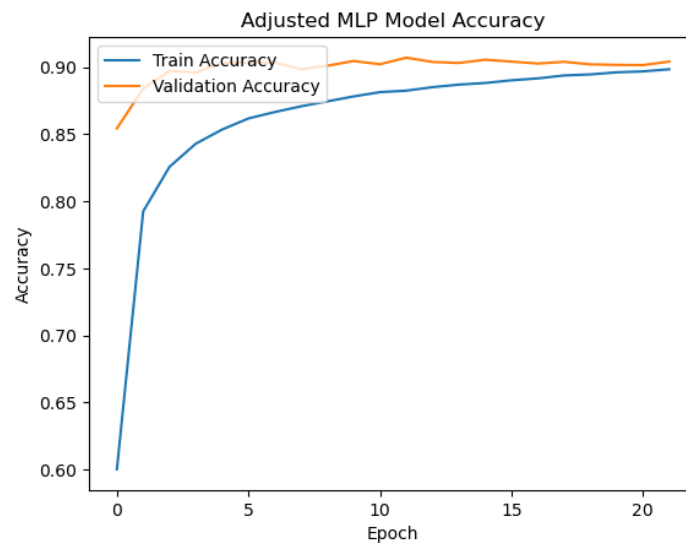


Figure 6: Adjusted MLP Model Accuracy




Figure 7: Adjusted MLP Model Loss

Table 4: Performance Metrics for Adjusted MLP on the Test Subset

Class	Precision	Recall	F1-Score	Support
ADJ	0.89	0.88	0.89	1788
ADP	0.90	0.96	0.93	2034
ADV	0.90	0.84	0.87	1176
AUX	0.95	0.99	0.97	1543
CCONJ	0.99	0.99	0.99	737
DET	0.98	0.98	0.98	1897
INTJ	0.00	0.00	0.00	120
NOUN	0.83	0.90	0.87	4137
NUM	0.77	0.65	0.70	542
PART	0.92	0.98	0.95	649
PRON	0.97	0.98	0.97	2162
PROPN	0.74	0.72	0.73	2077
PUNCT	0.97	0.99	0.98	3096
SCONJ	0.83	0.60	0.70	384
SYM	1.00	0.02	0.04	109
VERB	0.94	0.91	0.92	2606
X	0.00	0.00	0.00	39
-	0.97	0.97	0.97	354
Accuracy	0.90 (25,450)			
Macro Avg	0.81	0.74	0.75	25,450
Weighted Avg	0.90	0.90	0.90	25,450

796/796 ————— 1s 791us/step				
Adjusted MLP Test Accuracy: 0.9084872298624754				
Classification Report for the Adjusted MLP Model on Test Data:				
	precision	recall	f1-score	support
ADJ	0.91	0.88	0.89	1788
ADP	0.89	0.98	0.93	2034
ADV	0.91	0.84	0.87	1176
AUX	0.95	0.98	0.97	1543
CCONJ	0.98	0.99	0.99	737
DET	0.98	0.97	0.98	1897
INTJ	0.75	0.07	0.14	120
NOUN	0.81	0.93	0.86	4137
NUM	0.95	0.66	0.78	542
PART	0.95	0.96	0.95	649
PRON	0.94	0.98	0.96	2162
PROPN	0.80	0.70	0.75	2077
PUNCT	0.98	0.99	0.99	3096
SCONJ	0.95	0.55	0.70	384
SYM	0.94	0.15	0.25	109
VERB	0.93	0.94	0.94	2606
X	0.00	0.00	0.00	39
-	0.99	0.96	0.97	354
accuracy			0.91	25450
macro avg	0.87	0.75	0.77	25450
weighted avg	0.91	0.91	0.90	25450

Figure 8: Performance Metrics for Adjusted MLP on the Test Subset

8.9 Bi-Directional Stacked RNN

8.10 RNN Model Summary

The following table shows the architecture of the Bi-Directional Stacked RNN model. It includes an embedding layer to convert words into dense vectors, followed by bidirectional LSTM layers, dense, and dropout layers. The model uses a softmax activation in the output layer for multi-class classification.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 5, 300)	5,965,200
dropout_2 (Dropout)	(None, 5, 300)	0
bidirectional (Bidirectional)	(None, 5, 1024)	3,330,048
dropout_3 (Dropout)	(None, 5, 1024)	0
bidirectional_1 (Bidirectional)	(None, 1024)	6,295,552
dropout_4 (Dropout)	(None, 1024)	0
dense_3 (Dense)	(None, 64)	65,600
dropout_5 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 18)	1,170

Total params: 15,657,570 (59.73 MB)
Trainable params: 9,692,370 (36.97 MB)
Non-trainable params: 5,965,200 (22.76 MB)
None

Figure 9: Bi-Directional Stacked RNN Model Summary

8.11 Architecture

We implemented a bi-directional stacked RNN using LSTM cells. The architecture consisted of:

- **Embedding layer:** Converts words into 300-dimensional vectors using pre-trained Word2Vec embeddings.
- **Bidirectional LSTM layers:** Two bidirectional LSTM layers to capture dependencies in both directions.
- **Fully connected layer:** Dense layer with ReLU activation.
- **Dropout layers:** Prevents overfitting by randomly setting input units to zero.
- **Output layer:** Softmax activation for multi-class classification.

Figures 10 and Figure 11 depict the accuracy and loss of the Bi-Directional RNN model over 25 epochs for both training and validation datasets. The training accuracy steadily increases and converges around 0.85, while the validation accuracy closely follows, reaching similar values with minor fluctuations, indicating a well-generalized model. Figure 11 shows a sharp decline in both training and validation losses in the initial epochs, with the training loss gradually decreasing further and the validation loss stabilizing around 0.6. The close alignment of training and validation metrics suggests effective learning with minimal overfitting.




Figure 10: Bi-Directional RNN Model Accuracy




Figure 11: Bi-Directional RNN Model Loss

8.12 Training and Hyperparameter Tuning

The model was trained using the Adam optimizer. Early stopping and model checkpoints were used to monitor validation performance and prevent overfitting. Dropout layers with a rate of 0.5 were used to randomly set input units to zero during training.

8.13 Results

8.14 Evaluation Metrics

Table 5: Performance Metrics for Bi-Directional RNN on the Test Subset

Class	Precision	Recall	F1-Score	Support
ADJ	0.94	0.87	0.90	1788
ADP	0.85	0.89	0.87	2034
ADV	0.96	0.87	0.91	1176
AUX	0.99	0.99	0.99	1543
CCONJ	0.65	0.50	0.56	737
DET	0.90	0.94	0.92	1897
INTJ	0.80	0.85	0.82	120
NOUN	0.92	0.82	0.87	4137
NUM	0.75	0.54	0.63	542
PART	0.80	0.90	0.85	649
PRON	0.98	0.99	0.99	2162
PROPN	0.82	0.53	0.64	2077
PUNCT	0.59	0.90	0.71	3096
SCONJ	0.96	0.79	0.87	384
SYM	0.93	0.62	0.75	109
VERB	0.94	0.92	0.93	2606
X	0.00	0.00	0.00	39
-	0.88	0.89	0.88	354
Accuracy	0.85 (25,450)			
Macro Avg	0.81	0.77	0.78	25,450
Weighted Avg	0.86	0.85	0.85	25,450

The Bi-Directional Stacked RNN achieved a test accuracy of 0.85, demonstrating significant improvement over the baseline models.

796/796 ————— 10s 12ms/step
 Test Accuracy: 0.8877799607072692

Classification Report for the Test Set:

	precision	recall	f1-score	support
ADJ	0.94	0.93	0.93	1788
ADP	0.85	0.92	0.88	2034
ADV	0.97	0.89	0.93	1176
AUX	0.99	0.99	0.99	1543
CCONJ	0.76	0.46	0.57	737
DET	0.94	0.96	0.95	1897
INTJ	0.95	0.88	0.91	120
NOUN	0.91	0.91	0.91	4137
NUM	0.78	0.57	0.66	542
PART	0.85	0.93	0.89	649
PRON	0.99	0.99	0.99	2162
PROPN	0.81	0.67	0.73	2077
PUNCT	0.73	0.90	0.81	3096
SCONJ	0.91	0.85	0.88	384
SYM	0.93	0.62	0.75	109
VERB	0.95	0.97	0.96	2606
X	0.00	0.00	0.00	39
-	0.93	0.91	0.92	354
accuracy			0.89	25450
macro avg	0.84	0.80	0.81	25450
weighted avg	0.89	0.89	0.88	25450

Figure 12: Performance Metrics for bi-directional stacked RNN using LSTM cells

8.15 Precision-Recall AUC for Each Class

The Precision-Recall AUC provides a comprehensive view of the model's performance for each class, especially in the context of imbalanced data. The following figure illustrates the Precision-Recall curves and AUC scores for each class after running the Bi-Directional Stacked RNN model.

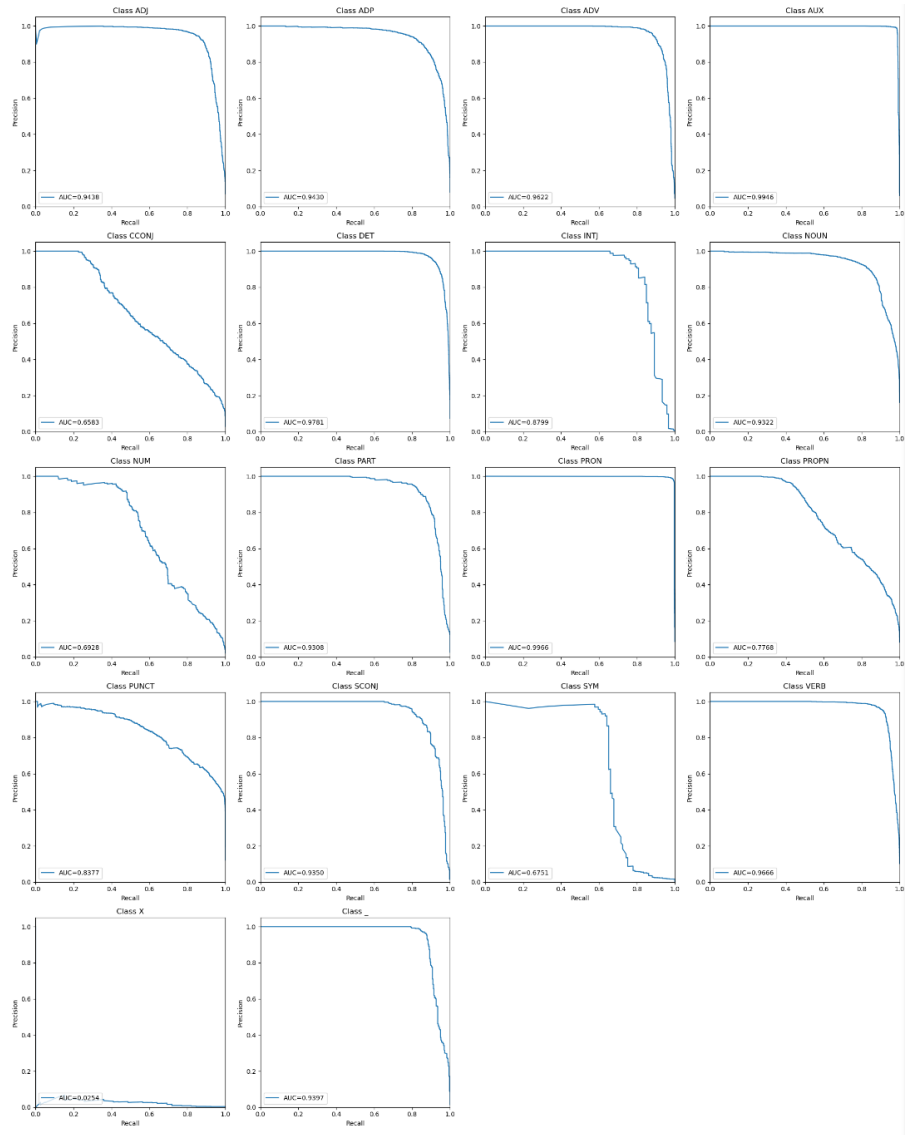


Figure 13: Precision-Recall AUC for Each Class

Table 6: Precision-Recall AUC for Each Class

Class	Precision-Recall AUC
ADJ	0.9438
ADP	0.9430
ADV	0.9622
AUX	0.9946
CCONJ	0.6583
DET	0.9781
INTJ	0.8799
NOUN	0.9322
NUM	0.6928
PART	0.9308
PRON	0.9966
PROPN	0.7768
PUNCT	0.8377
SCONJ	0.9350
SYM	0.6751
VERB	0.9666
X	0.0254
-	0.9397

8.16 Insights from Precision-Recall AUC

The Precision-Recall AUC scores reveal the following insights:

- Classes such as **AUX**, **PRON**, and **VERB** achieved high AUC scores, indicating the model’s strong ability to distinguish these classes effectively.
- The **CCONJ** and **SYM** classes had lower AUC scores, suggesting that these classes are more challenging for the model to predict accurately.
- The **X** class had a notably low AUC score, highlighting the model’s difficulty in identifying this rare class, which could be attributed to insufficient representation in the training data.
- Most major classes like **NOUN**, **ADJ**, **ADP**, and **DET** achieved high AUC scores, demonstrating the model’s overall strong performance across common POS tags.

8.17 Insights and Discussion

The Bi-Directional Stacked RNN demonstrated significant improvement over both baseline models. The MLP achieved notable performance gains; however, the RNN showed further enhancements across most metrics, except for a few classes like CCONJ and NUM. The learning curves indicate that the RNN effectively learned from the training data, although there were signs of slight overfitting in some classes.

8.18 Hyperparameter Tuning Insights

The hyperparameter tuning process revealed the following optimal settings for the Bi-Directional RNN:

- **Embedding dimension:** 300
- **LSTM units:** 128
- **Dropout rate:** 0.5
- **Batch size:** 64
- **Learning rate:** 0.001

8.19 Impact of Data Preprocessing

Extensive preprocessing steps, including text normalization, tokenization, and lemmatization, contributed to the improved performance of the RNN classifier. By removing noise and standardizing the text data, the model could focus on the most relevant features for POS classification.

8.20 Conclusion

The Bi-Directional Stacked RNN for POS tagging demonstrated superior performance compared to baseline models, achieving a test accuracy of 0.85. The model effectively learned to classify POS tags, leveraging both past and future context in the sequences. Future work could explore more advanced architectures and larger datasets to further enhance performance. Additionally, incorporating character-level embeddings and experimenting with other regularization techniques could yield even better results.